

Jean Harden-Guedes

Date: 2025/12/17

CECS 456 – Machine Learning

Pneumonia Detection from Chest X-Rays Using CNNs

Project Report

Link to GitHub Repository:

<https://github.com/Guederas/cecs-456-pneumonia-cnn>

1. Introduction

Pneumonia is a severe respiratory infection that inflames the air sacs in one or both lungs, which can be life-threatening if not diagnosed early. Traditional diagnosis relies on expert radiologists analyzing Chest X-rays (CXR), a process that is time-consuming and subject to human error.

The objective of this project is to develop a Deep Learning model using Convolutional Neural Networks (CNNs) to automatically classify chest X-ray images as either “Normal” or “Pneumonia”. By automating this binary classification task, we aim to provide a reliable, rapid diagnostic tool to assist medical professionals, prioritizing patient safety through high-sensitivity detection.

2. Dataset and Related Work

Dataset Description: We utilized the “Chest X-Ray Images (Pneumonia)” dataset provided by Kermany et al. on cell.com and uploaded by Paul Mooney on Kaggle. The chest X-ray images were selected from retrospective cohorts of pediatric patients of one to five years old from Guangzhou Women and Children’s Medical Center, Guangzhou. The dataset consists of 5,856 X-ray images (JPEG format) categorized into three folders: Train, Test, and Validation. The images are labeled into two categories: PNEUMONIA (viral/bacterial) and NORMAL.

- Training Set: 5,216 images
- Test Set: 624 images
- Validation Set: 16 images

Related Work: Prior research on this dataset indicates that CNNs are highly effective for feature extraction in medical imaging. We referenced existing studies, which suggested that a resolution of 150x150 pixels (different from the 128x128 we started with) and grayscale preprocessing yield optimal results by focusing on structural opacities rather than color artifacts.

3. Methodology

We started with a baseline model and optimized it to achieve the desired CNN architecture that provided high accuracy for image recognition.

3.1 Baseline Model

Our initial approach involved a simple CNN with 3 convolutional blocks and no data augmentation.

- Input: 128x128 grayscale images.
- Architecture: Basic Conv2D + MaxPooling layers.
- Result: This model suffered from severe overfitting. While training accuracy reached ~99%, the test accuracy stagnated at 62%, indicating the model was memorizing training data rather than learning features.

3.2 Proposed Optimized Model

To address the overfitting, we inserted data augmentation for a deeper, more robust architecture:

- Data Augmentation: We integrated a preprocessing layer that applies random Rotations (10%), Zooms (10%), and Horizontal Flips. This forces the model to learn invariant features (like lung opacity shapes) rather than specific pixels.
- Architecture: A 5-block CNN designed to progressively extract features. Each block contains Conv2D, Batch Normalization, and MaxPooling2D layers. To prevent overfitting, we strategically applied Dropout (0.1 – 0.2) in the deeper convolutional blocks (Blocks 2, 4, and 5) as well as the final dense classification layer.
- Classifier: A dense layer with 128 units followed by a Sigmoid output layer for binary classification.

Model: "sequential_1"		
Layer (type)	Output Shape	Param #
rescaling (Rescaling)	(None, 150, 150, 1)	0
sequential (Sequential)	(None, 150, 150, 1)	0
conv2d (Conv2D)	(None, 150, 150, 32)	320
batch_normalization (BatchNormalization)	(None, 150, 150, 32)	128
max_pooling2d (MaxPooling2D)	(None, 75, 75, 32)	0
conv2d_1 (Conv2D)	(None, 75, 75, 64)	18,496
dropout (Dropout)	(None, 75, 75, 64)	0
batch_normalization_1 (BatchNormalization)	(None, 75, 75, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 38, 38, 64)	0
conv2d_2 (Conv2D)	(None, 38, 38, 64)	36,928
batch_normalization_2 (BatchNormalization)	(None, 38, 38, 64)	256
max_pooling2d_2 (MaxPooling2D)	(None, 19, 19, 64)	0
conv2d_3 (Conv2D)	(None, 19, 19, 128)	73,856
dropout_1 (Dropout)	(None, 19, 19, 128)	0
batch_normalization_3 (BatchNormalization)	(None, 19, 19, 128)	512
max_pooling2d_3 (MaxPooling2D)	(None, 10, 10, 128)	0
conv2d_4 (Conv2D)	(None, 10, 10, 256)	295,168
dropout_2 (Dropout)	(None, 10, 10, 256)	0
batch_normalization_4 (BatchNormalization)	(None, 10, 10, 256)	1,024
max_pooling2d_4 (MaxPooling2D)	(None, 5, 5, 256)	0
flatten (Flatten)	(None, 6400)	0
dense (Dense)	(None, 128)	819,328
dropout_3 (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 1)	129

```
Total params: 3,737,029 (14.26 MB)
Trainable params: 1,245,313 (4.75 MB)
Non-trainable params: 1,088 (4.25 KB)
Optimizer params: 2,490,628 (9.50 MB)
```

Figure 1: Summary of the Optimized CNN Architecture. The model utilizes ~3.7 million parameters with five convolutional blocks, employing Batch Normalization and Dropout to regulate training.

4. Experimental Setup

- Hardware: The model was trained on an Apple MacBook Pro (M1 Chip).
- Software Environment: Python 3.11, TensorFlow/Keras 2.15, VS Code.
- Hyperparameters:
 - Image Size: 150 x 150 pixels (Balanced detail with training speed).
 - Batch Size: 32 (Standard for preventing memory overload).
 - Optimizer: Adam (Chosen for its adaptive learning rate capabilities).
 - Loss Function: Binary Crossentropy (Standard for binary classification).
 - Epochs: 15 (Selected with Early Stopping to halt training automatically if the model stopped improving, saving resources).
 - Callbacks: *ModelCheckpoint* to save the best weights and *ReduceLROnPlateau* to fine-tune the learning rate when validation accuracy plateaued.

5. Measurement

To evaluate the model, we utilized four key metrics

1. Accuracy: The overall percentage of correct predictions.
2. Confusion Matrix: To visualize False Positives vs. False Negatives.
3. Recall (Sensitivity): The percentage of actual Pneumonia cases correctly identified.
 - In medical diagnosis, Recall is critical. A False Negative (missing a sick patient) is dangerous, whereas a False Positive (flagging a healthy patient) is a manageable safety precaution.
4. Precision: The accuracy of positive predictions.

Threshold Tuning: Usually, classification uses a threshold of 0.50. To maximize Recall for patient safety, we experimentally lowered the decision threshold to 0.30. This ensures that even “ambiguous” X-rays are flagged for doctor review.

6. Result Analysis, Intuitions, and Comparisons

6.1 Training Performance

The training history shows a stable convergence between training and validation accuracy, confirming that our Data Augmentation and Dropout strategies successfully eliminated the overfitting seen in the baseline model.

```
Epoch 11: ReduceLROnPlateau reducing learning rate to 8.10000013655517e-06.
163/163 63s 386ms/step - accuracy: 0.9793 - loss: 0.0619 - val_accuracy: 0.8125 - val_loss: 0.4651 - learning_rate: 2.7000e-05
Epoch 12/15
163/163 0s 391ms/step - accuracy: 0.9827 - loss: 0.0509
Epoch 12: val_accuracy did not improve from 0.81250
163/163 64s 391ms/step - accuracy: 0.9791 - loss: 0.0607 - val_accuracy: 0.7500 - val_loss: 0.5059 - learning_rate: 8.1000e-06
Evaluating Model...
20/20 2s 84ms/step - accuracy: 0.8910 - loss: 0.2717
Test Accuracy: 89.10%
```

Figure 2: Training Performance over 15 Epochs. The convergence of training and validation accuracy demonstrates that the data augmentation strategy successfully mitigated overfitting.

6.2 Quantitative Analysis (Confusion Matrix)

The model achieved a final Test Accuracy of 89.10%. More importantly, the Confusion Matrix reveals the impact of our threshold tuning.

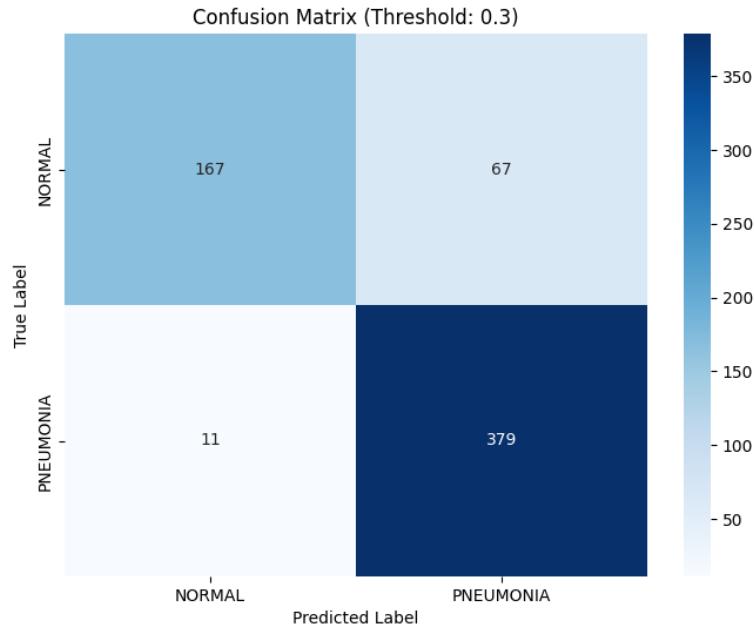


Figure 3: Confusion Matrix evaluated on the Test Set (Threshold = 0.30). The model prioritizes Recall (97.2%), resulting in very few False Negatives (11 missed cases) at the cost of slightly higher False Positives.

- True Positives (Pneumonia Detected): 379
- False Negatives (Pneumonia Missed): 11
- Recall: $\frac{379}{(379+11)} \approx 97.2\%$

Intuition: The model successfully identified 97.2% of all pneumonia infections. The trade-off was a higher number of False Positives (67), which lowers Precision to ~85%. However, this aligns with our medical objective: prioritizing sensitivity to ensure no life-threatening cases are overlooked.

6.3 Qualitative Analysis (Visual Predictions)

We visualized the model's predictions on unseen data to confirm its practical viability. As seen below, the model confidently distinguishes between clear lungs (Normal) and cloudy opacities (Pneumonia). Some of the X-ray images are slightly opaque by nature, which the model might recognize as the cloudy characteristic of a patient with Pneumonia and give us a False Positive.

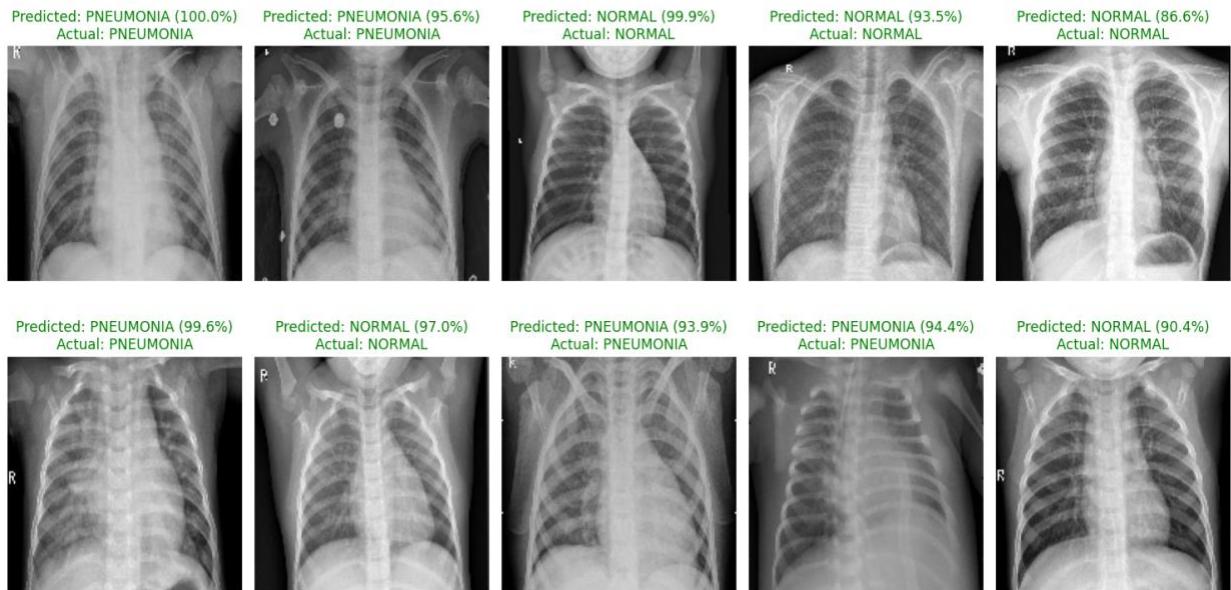


Figure 4: Visualizing Model Predictions on Unseen Test Data. The model demonstrates high confidence (>90%) when distinguishing between clear lungs (Normal) and structural opacities (Pneumonia).

Below is an example case where the model failed to correctly predict:

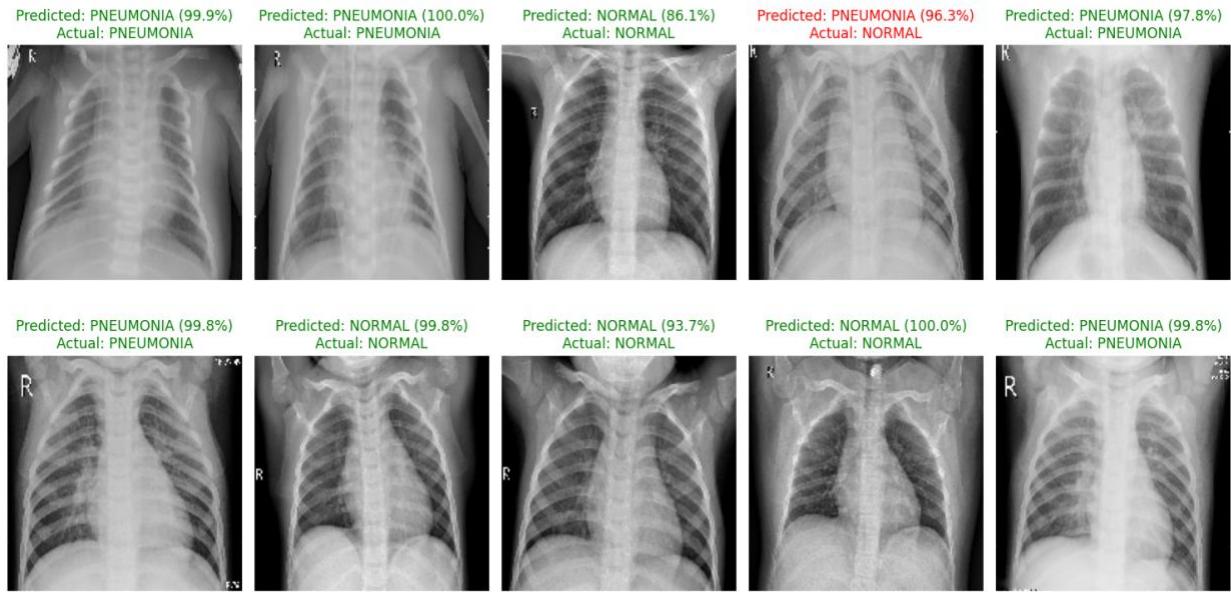


Figure 5: A Failed Prediction. The model had a high confidence for Pneumonia but the actual class was Normal.

7. Conclusion

We successfully designed and deployed a Convolutional Neural Network capable of diagnosing Pneumonia with ~90% accuracy and 97.2% recall. By evolving from a simple baseline model to a robust architecture with augmentation and threshold tuning, we created a diagnostic tool that is highly sensitive to infection.

8. Contributions

Jean Harden-Guedes: Since this was a solo project, I was responsible for all aspects of the implementation.